

lab

January 21, 2021

1 Week 6. Optimization. Programming Task

1.0.1 For grading

```
[1]: #DO NOT CHANGE
import grading
import grading_utils
```

```
[2]: #DO NOT CHANGE
grader = grading.Grader(assignment_key="Aoii7s9WRQyvONf6kwwjAw",
                        all_parts=["NCrTc", "JigtH", "HlACv"])
```

```
[18]: # token expires every 30 min
COURSERA_TOKEN = "0AoIh4QOahl5nZne"### YOUR TOKEN HERE
COURSERA_EMAIL = "manuelalejandromartinezf@gmail.com"### YOUR EMAIL HERE
```

Let us consider the **House Pricing** dataset, where you have a lot of information about the houses being sold and you aim to produce the price of the house.

Firstly, let us import basic libraries ([numpy \(docs\)](#) for matrix operations and [pandas \(docs\)](#) for convinient dataset workaround):

```
[4]: import numpy as np
import pandas as pd
```

1.0.2 Task 1. Reading and Preparing

```
[5]: datX=np.load('x_train.npy')
datY=np.log(np.load('y_train.npy'))
datX=pd.DataFrame(datX, columns=datX.dtype.names)
datX
```

```
[5]:
```

	date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
0	2014-09-16	5.0	3.25	3710	34200	2.0	
1	2014-11-18	3.0	1.75	2820	8879	1.0	
2	2014-11-10	3.0	1.00	1240	239144	1.0	

3	2015-04-16	4.0	2.50	2670	8279	2.0
4	2014-07-23	3.0	2.25	2700	4025	2.0
...
14995	2014-05-21	4.0	2.75	2290	6120	2.0
14996	2015-04-01	3.0	2.00	1430	9250	1.0
14997	2014-07-11	2.0	1.00	640	7768	1.0
14998	2014-05-15	3.0	1.00	1630	10304	1.0
14999	2014-11-20	2.0	1.00	720	4592	1.0

	waterfront	condition	grade	sqft_above	sqft_basement	yr_built	\
0	False	3	8	2510	1200	1986	
1	False	5	7	1540	1280	1920	
2	False	3	6	1240	0	1921	
3	False	3	7	2670	0	1999	
4	False	4	8	1760	940	1907	
...	
14995	False	4	7	2170	120	1926	
14996	False	4	8	990	440	1983	
14997	False	3	6	640	0	1942	
14998	False	5	7	1630	0	1953	
14999	False	4	6	720	0	1943	

	yr_renovated	zipcode	lat	long
0	0	98074	47.610100	-122.046997
1	1957	98146	47.509399	-122.375999
2	1992	98038	47.430302	-122.045998
3	0	98148	47.429199	-122.328003
4	0	98122	47.607399	-122.293999
...
14995	0	98115	47.674599	-122.327003
14996	0	98052	47.695202	-122.096001
14997	0	98106	47.514999	-122.359001
14998	0	98155	47.754799	-122.317001
14999	0	98199	47.653400	-122.403999

[15000 rows x 16 columns]

Okay, we manage to load the data (you can read more about the [load here](#). But it is not a necessity). We are going to use linear models to work with it, but firstly we need to come up with idea what features should we include in the model at all (which feature the price is lineary dependent on):

Do not forget to install seaborn. You can do that by running `pip install seaborn` in the command line locally, or simply by running the next cell:

```
[6]: !pip install seaborn
```

Requirement already satisfied: seaborn in /opt/conda/lib/python3.7/site-packages

(0.10.1)
 Requirement already satisfied: pandas>=0.22.0 in /opt/conda/lib/python3.7/site-packages (from seaborn) (1.0.3)
 Requirement already satisfied: numpy>=1.13.3 in /opt/conda/lib/python3.7/site-packages (from seaborn) (1.18.4)
 Requirement already satisfied: matplotlib>=2.1.2 in /opt/conda/lib/python3.7/site-packages (from seaborn) (3.2.1)
 Requirement already satisfied: scipy>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from seaborn) (1.4.1)
 Requirement already satisfied: pytz>=2017.2 in /opt/conda/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn) (2020.1)
 Requirement already satisfied: python-dateutil>=2.6.1 in /opt/conda/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn) (2.8.1)
 Requirement already satisfied: cycloer>=0.10 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn) (0.10.0)
 Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn) (1.2.0)
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn) (2.4.7)
 Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas>=0.22.0->seaborn) (1.14.0)

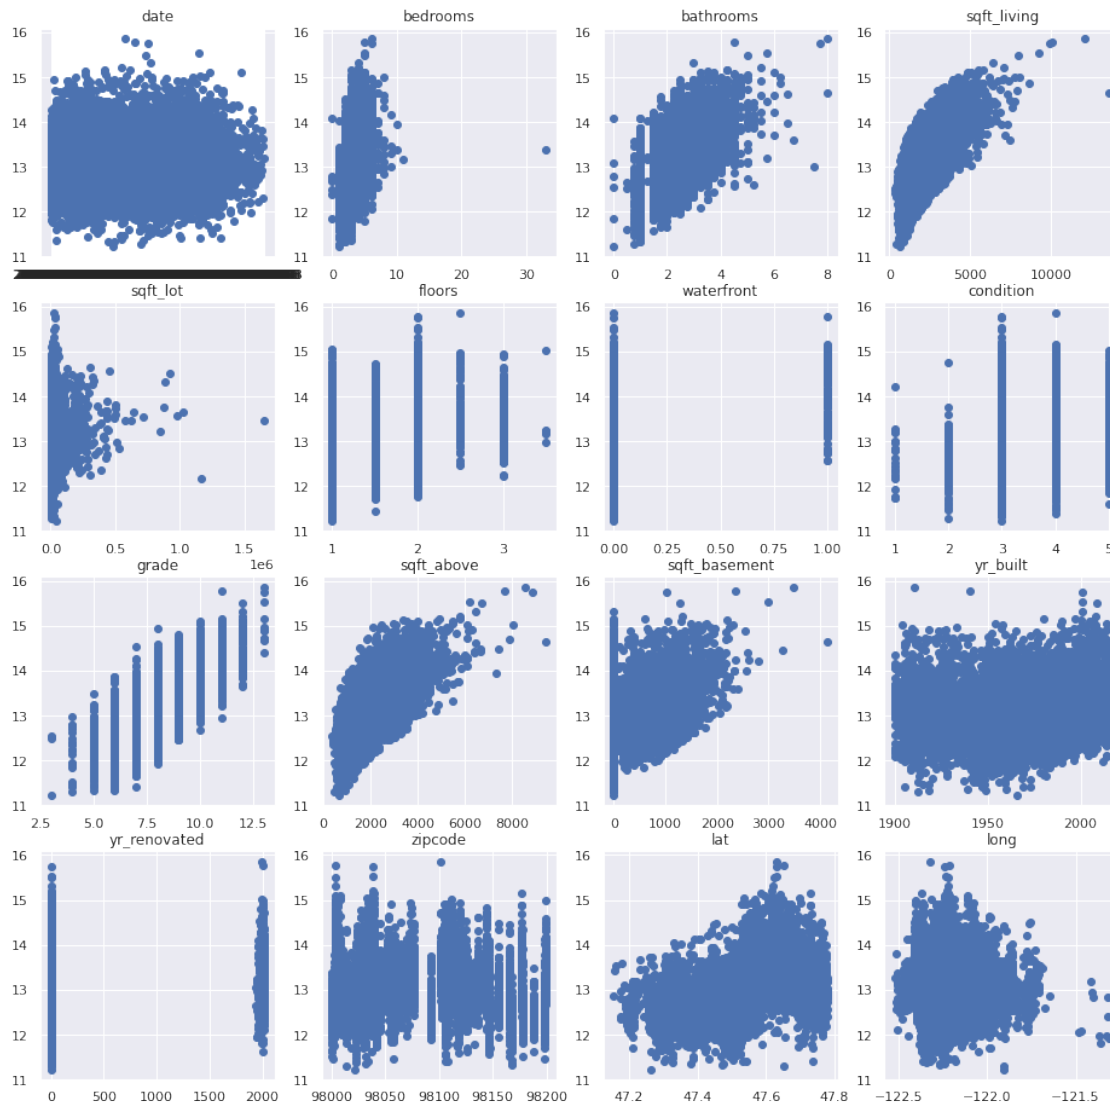
In order to do it let us plot every feature vs the price. Firstly, we import nice plotting modules:

```
[7]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

%matplotlib inline
```

```
[8]: f, ax=plt.subplots(4, 4, figsize=(16,16))

for i, name in enumerate(datX.columns):
    ax[i//4][i%4].scatter(datX[name], datY)
    ax[i//4][i%4].set_title(name)
```



Let us say, that we choose to work the following set of features: + bedrooms + bathrooms + sqft_living + floors + condition + grade + sqft_above + sqft_basement + long + lat

Clear the dataset from all the other features and create: 1. matrix X , all elements should be real numbers 2. number N – number of considered houses 3. number m – number of new features

Hint: it is easier to clean columns from dataset (you should look [here](#) for inspiration) and then get a matrix with `.values`

Warning: Please use features in the order mentioned above!

[74]:

```
#your code goes here
X=(datX[["bedrooms","bathrooms","sqft_living","floors","condition","grade","sqft_above","sqft_basement","long","lat"]].values)
N=np.shape(X)[0]
```

```
m=np.shape(X)[1]
```

Run the following cells to automatically check results of your code:

```
[10]: ## GRADED PART, DO NOT CHANGE!
grader.set_answer("NCrTc", grading_utils.test_reader(X, N, m))
```

```
[11]: # you can make submission with answers so far to check yourself at this stage
grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

You used an invalid email or your token may have expired. Please make sure you have entered all fields correctly. Try generating a new token if the issue still persists.

Consider that we are interested in the loss of the model we discussed in the video:

- Assume we have input data that is denoted as $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_N$
- House prices for this input data are known y_1, y_2, \dots, y_N

We propose a **simple linear model** for this task:

$$\hat{y}_i = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$$

As a loss function we will use the mean squared error (**MSE**):

$$Loss(\vec{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

1.0.3 Task 2. Compute analytically the $Loss(\vec{w})$ function

Please, keep the signature of the function and enter the code only under **your code goes here**. **Attention:** try to avoid usage of **for** cycles! The easiest way to do it is by using matrix operations.

Hint: to get nice w_0 coefficient it is convenient to add to the **X** matrix the column of 1 with `np.concatenate` [documentation](#)

```
[12]: def loss(w, X, y):
      #your code goes here
      X=np.concatenate((np.ones(shape=(N,1)),X),axis=1)
      lossValue=np.mean((y-(np.dot(X,w.T)))**2)
      return lossValue
```

Run the following cells to automatically check results of your code:

```
[13]: ## GRADED PART, DO NOT CHANGE!
grader.set_answer("JigtH", grading_utils.test_loss(loss, X, datY))
```

```
[14]: # you can make submission with answers so far to check yourself at this stage
grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

You used an invalid email or your token may have expired. Please make sure you have entered all fields correctly. Try generating a new token if the issue still persists.

1.0.4 Task 3. Compute analytically the gradient of the $Loss(\vec{w})$

Please, enter your answer in the cell below (it should be a `markdown` cell). You can either specify each partial derivative $\frac{\partial Loss}{\partial w_i}$ or $\nabla Loss$ altogether using matrix operations.

$$\frac{\partial}{\partial w} = -2/N \sum (y_i - w_0 - w_1 x_i) x_i$$

1.0.5 Task 4. Write a function to compute the gradient of the Loss function in the given point

Please, keep the signature of the function and enter the code only under `your code goes here`. **Attention:** try to avoid usage of `for` cycles! The easiest way to do it is by using matrix operations.

```
[42]: def grad(w_k, X, y):
      #your code goes here
      N = X.shape[0]
      one = np.ones((X.shape[0], 1))
      Xbar = np.concatenate((one, X), axis = 1)
      return 2/N * (Xbar.T).dot(Xbar.dot(w_k)-y)
```

Run the following cells to automatically check your function.

```
[20]: ## GRADED PART, DO NOT CHANGE!
grader.set_answer("H1ACv", grading_utils.test_grad(grad, X, datY))
```

```
[21]: # you can make submission with answers so far to check yourself at this stage
grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

1.0.6 Task 5. Write gradient descent

Now it is time to formulate the gradient descent! As you remember, the idea here is that:

$$\vec{w}^{k+1} = \vec{w}^k - \alpha_k \cdot \nabla Loss(\vec{w}^k)$$

We propose that you use constant $\alpha_k = \alpha$. Assume that the method should stop in two cases: + if the number of iterations is too high (`maxiter`) + if the length of the gradient is low enough (`<eps`) to call an extremum

Please, keep the signature of the function and enter the code only under `your code goes here`.

```
[112]: def gradDescent(w_init, alpha, X, y, maxiter=500, eps=1e-2):
    losses=[]
    w=[w_init]

    for it in range(maxiter):
        w_new = w[-1] - alpha*grad(w[-1],X,y)
        if np.linalg.norm(grad(w_new,X,y))/len(w_new) < eps:
            break
        w.append(w_new)
        losses.append(loss(w_new,X,y))
    return losses,w[-1], it
```


Experiment with several alphas and several initial values of weights. To illustrate, provide graphs for the Loss function over iterations in each case (and, optionally, the distance between weights from one iteration to the next):

(we provided all key plotting commands for you, but you can always look into [this tutorial](#))

```
[121]: plt.figure(figsize=(8,8))

#code goes here
w = np.array([1,1,1,1,1,1,1,1,1,1,1])
losses, w,it = gradDescent(w.T,0.01, X, datY,5)
print(losses)
plt.plot( np.arange(it+1),losses, 'r')
plt.ylabel('losses')
plt.xlabel('iterations (per hundreds)')
plt.title("Learning rate = 0.01")
plt.show()
```

```
[6.760047768595087e+17, 2.264243682946674e+28, 7.584008047014016e+38,
2.540237994948272e+49, 8.508441751388078e+59]
```



Let us check the adequacy of the model we created.

Choose several (no less than five) houses (inputs in your **X** matrix) and calculate predicted prices by:

$$\hat{y}_i = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$$

```
[122]: test = X[:5]
one = np.ones((test.shape[0], 1))
Xbar = np.concatenate((one, test), axis = 1)

y_hat = Xbar.dot(w)
y_hat = np.array(y_hat)
```



```
np.all((y_hat-datY[:5] < 0.001))
```

[122]: True

Compare predicted values with an actual answer (stored in your y array). Is it satisfying enough?

1.0.7 Task 6. Discussion

Answer following questions: 1. Does your method converge at least for some alpha? If not, what might be the workaround? 2. How does changing of the alpha influence the speed of convergence? 3. Are the optimal weights in all convergent cases the same? 4. How does this affect the Loss function?